

Data storage

- ① HDFS
- ② HBase

• Data processing & Analytics

- Batch
 - map reduce
 - Spark
 - Hive

• Online streaming process

- storm
- Spark

• Data sharding

- Data partitioned into groups & stored on different machines

ex. HDFS

NO SQL DBDS

- Concurrent processing of data on different machines
- Speed up processing time.

- Distributed Algorithms

Set of collaborative tasks which might run concurrently to process data on different machines.


Example -


- Hadoop Map reduce
- Distributed ML

framework to execute

- Hadoop YARN

HDFS





HDFS uses data sharding to create blocks of data & then it is distributed

data co-ordinator = Namenode

shard manager = datanode.

- Namenode Responsibility

1. - metadata

- tree structure

- directory metadata

- files metadata

2. location of each block

stores in this way

- Namespace img files

fsimage - Latest checkpoint

Datanode Responsibility

- preserve the assigned blocks on its local disk
- Each block is identified using blockpool id & block id
- Management of assigned blocks.

HDFS Availability

- large cluster = High chances of failure
failure = datanode crash, datablock on that datanode is not available
- If namenode crash HDFS unavailable for client

Two types of redundancy

- ① Block
- ② Namenode

① Block -

- Replicas are persisted on different m/c
- Number of block replicas are configurable with HDFS-site.xml
- client can read from any copy
- client write to all copies

② Namenode Redundancy


① Additional secondary namenode can be implemented

primary crashed secondary takes over

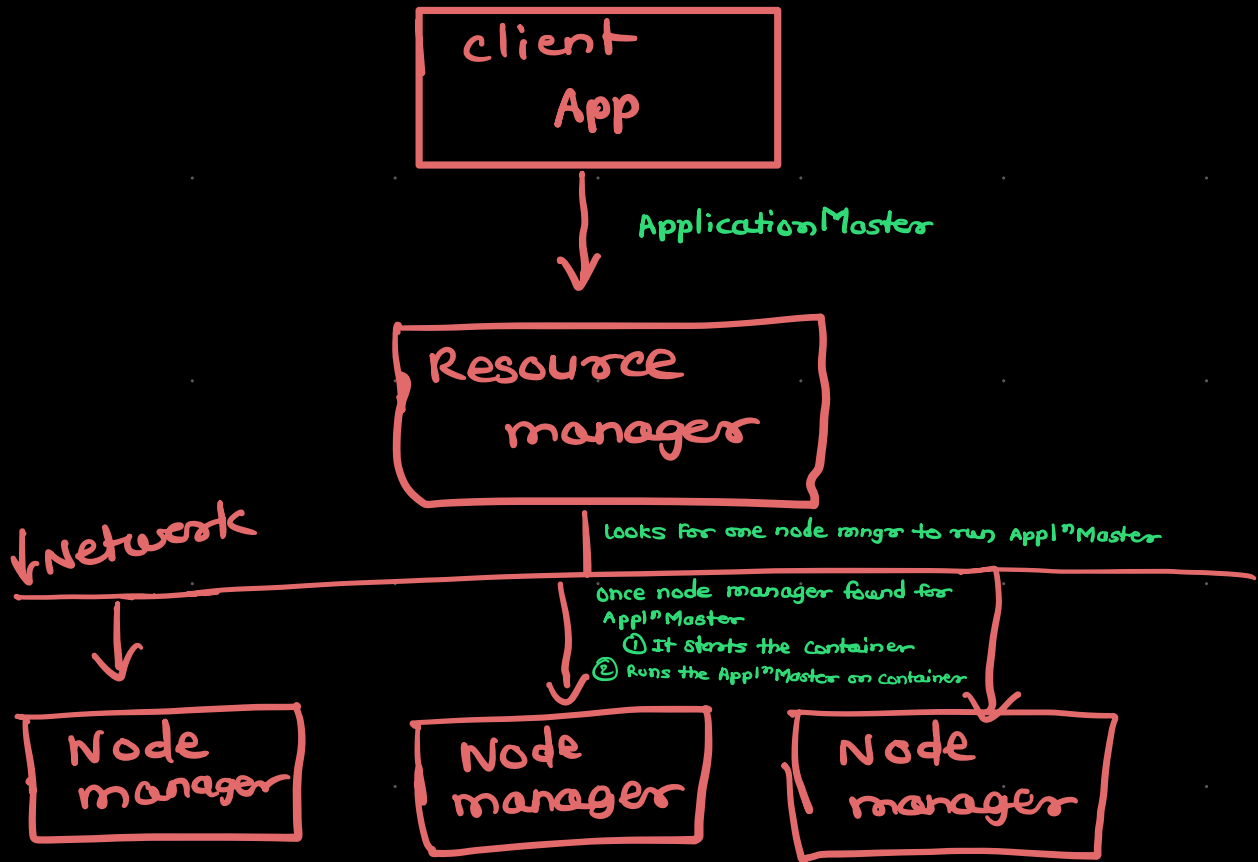
Two types of namenode fail-over

① Pre-release 2 namenode fail

② Release 2 namenode failovers.



Hadloop YARN



- AS YARN uses distributed algorithm It is implemented as YARN Applⁿ Master.


client job sends ApplicationMaster to → Resource manager

↓
look for node managers to run Applⁿ manager
once node manager selected, resource manager will ask the node manager to

- ① start container
- ② Run ApplicationMaster on container

* YARN Job Scheduling

- ① FIFO scheduler
- ② Capacity scheduler
- ③ fair scheduler



- perform aggregate operations on each group

- join two files with a given join key.

Two phases

① Map phase

— Retrieve data from local
redistribute the local data

② Reduce phase

Apply analytical logic to the data coming from other

mapper
→ combine →
to aggregate

reducer

