

- **ETCD clusters** - used to keep track of deployments and loggings.
- **kube-scheduler** - To schedule on dockers and which containers to use.
 - doesn't actually do it though
- **Controllers** -
 - Node controller
 - Replication-controller

• To monitor status & remediate situations

- **kube-apiserver** - Manages & controls all the services on the cluster.

- **kubelet** - Is an agent which runs on each node and act as agent.
It runs on that pods and monitor it basically acting as a captain of the ship.

- **kube-proxy** - To make sure proper rules are applied on worker-nodes to allow them to communicate.
 - Always look for a new service once found create appropriate rules and allow communication among pods that is done using creating IPTable.

ContainerD

① ctr - CLI for debugging

② nerdctl - alternative to docker

* crictl - used mainly for debugging purposes.

[from kubernetes community]

1.1 ETCD -

its a database

- key:value store database.
- By default when installed runs on port 2379 ← its a place where etcd listens
- By default you get 'etcdctl' ← client

• /etcdctl --version

• etcd runs as a pod in k8 which can be accessed using
kubectl get pods -- kube-system
• etcd-service

ETCD Commands

	vs
• etcd backup	• etcd snapshot save
• → cluster-health	• endpoint health
• → mk	• get
• → mkdir	• put
• → set	

ETCD

- stores information regarding cluster

ex. - Nodes - Accounts
 - PODS - Roles
 - Configs - Bindings
 - Secrets - others

- Two types of deployments

- ① from scratch

 ② from kubeadm tool

1.2 kubeAPI - Server

- `kubectl get pods -n kube-system`

gives all the pods used for kube system.

Location - /etc/kubernetes/manifest/kube-apiserver.yaml

- kube API lets end users, different parts of cluster and external components communicate with one another.
- Responsible for all orchestration with the servers that's why kube-apiserver.service consist of different certificates along with different services and their location.
- kubeAPI Server is at the center for all command
 - Authenticate User
 - validate Request
 - Retrive data
 - Update ETCD
 - Scheduler
 - kublet

location cat /etc/kubernetes/manifest/kube-apiserver.yaml.

View api - Server Options

- `ps -aux | grep kube-apiserver`

1.3 kube controller

controller-manager -

Is a process which continuously manages a components works towards maintaining desired status

- Watch status - auto heal
- Remediate situation.

① Node controller -

Monitor the state and take action's necessary to do so.

- There are number of controllers once installed they are located at
 - kube-controller-manager.service

↑

We have in .service

-- node monitoring-period = 5 sec
so once the node container is not responding it will go into grace period which is set to be 40 sec if container still doesn't reply then eviction timeout of 5m after that new container deployed.

This settings and much more can be changed at kube-controller-manager.service.

- By default all controllers are enabled.

- kubeadm

- kubectl get pods -n kube-system

```
cat /etc/kubernetes/manifests/kube-controller-manager.yaml
```

1.4 kube scheduler

- To schedule which pod goes on which container.

How does it decides ?

- Based on resource requirement

for ex. we need mem of 10gb in that case it will choose based on criteria to which containers has 10gb mem available while considering other factors such as CPU usage.

- - ① filter nodes
 - ② rank nodes

1.5 kubelet — captain on the ship

- ① Register a node
- ② Create a pods
- ③ Monitor Node 4 pods
- kubeadm does not deploy kubelets
Has to be manually installed

1.6 kube-proxy —

objective of kube-proxy -

- ① Look for new pods
- ② Once new pod found create a set of rules on each node to forward traffic b/w each node
- creates an IP table on each node to forward traffic

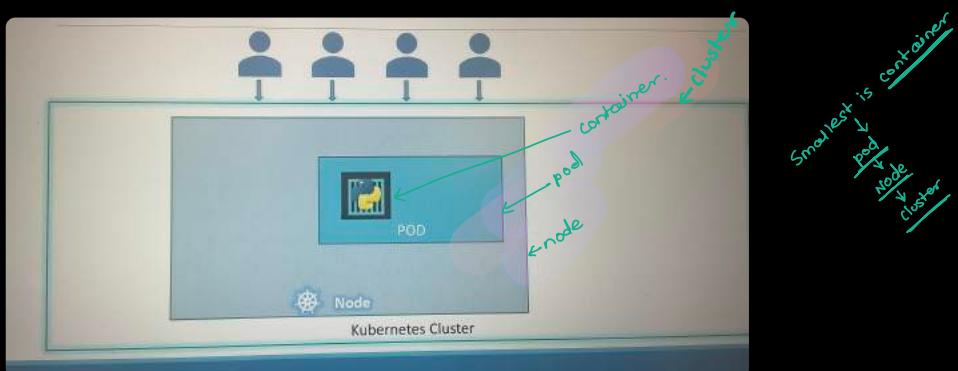
1.7 pods

- kubectl run nginx --image nginx
- kubectl get pods
- YAML files works as input for kubernetes to run applications
- kubectl describe pod myapp-pod

Yaml to create an pod contains

```
apiVersion: v1
kind: Pod
metadata:
spec:
```

AKMS



• Controllers

• Replication Controllers

- ① To load Balancing and Scaling
- ② To maintain the set replication pods running

- **Replica set** is a new & recommended way to set the replication set

rc - definition.yaml

• kubectl create -f rc - definition.yaml

• kubectl get replicationcontrollers
• k get replicaset
* New

replicaset - definition.yaml

difference between replication set and replica set is

- **replicaset** does require selector in definitions file.
- * Selectors are mentioned here so that replicaset can manage other pods mentioned in selector manually which matches the name explained in selector

Selector can also be used in replication set but they are optional.

To replace/update the definition file changes.

```
kubectl replace -f replicaset-definition.yaml
```

option # 2

- kubectl scale --replicas=6 -f

```
replicaset-definition.yaml
```

- kubectl get replicsets

- kubectl delete replicaset myapp-replicaset

- kubectl replace -f replicaset-def.yaml

- kubectl edit replicaset new-replica-set
 ↑
 name

Scaling

- kubectl scale rs new-replica-set --replicas=5

→ Replicaset definition file requires a selectors because it can also manage previously created pods.

• rc - definition.yaml

```
rc-definition.yaml
apiVersion: v1
kind: ReplicationController
metadata:
  name: myapp-rc
  labels:
    app: myapp
    type: front-end
spec:
  template:
    metadata:
      name: myapp-pod
      labels:
        app: myapp
        type: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 3
  selector: {no selectors}
```

this is the info used to create a replication controller

will be same because no selector

• replicaset - definition.yaml

```
replicaset-definition.yaml
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: myapp-replicaset
  labels:
    app: myapp
    type: front-end
spec:
  template:
    metadata:
      name: myapp-pod
      labels:
        app: myapp
        type: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 3
  selector: {selectors present}
```

matchLabels: {type: front-end}

replicaset created

replicaset status

replicaset pods

Name	Status	Reason	Pods
myapp-replicaset-0	Ready		1/1
myapp-replicaset-1	Running		2/2
myapp-replicaset-2	Running		2/2



Replica sets
Deployments
definition file for both will be same except the kind will be diff

• Deployments

- Deployments are used to manage the Replicasets.

- Rolling update is one of the main feature in deployment

- kubectl get all

- k get deployments

Services → enables connectivity betn pods

↳ This is how networking in k8 is done

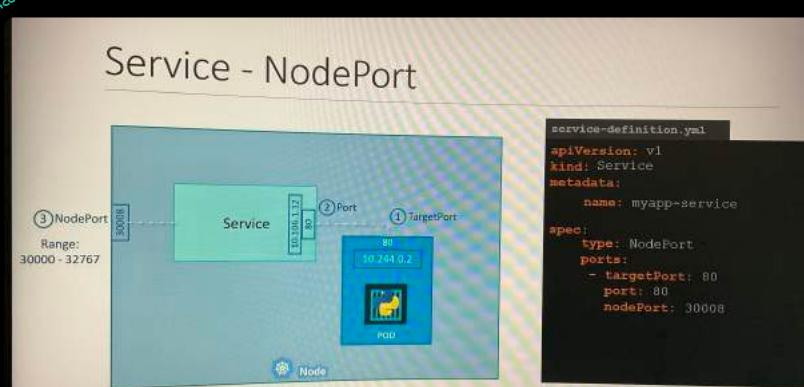
- Service types

- ① Node port
- ② clusterIP
- ③ Load Balancer

① Nodeport -

range - 30,000 - 32,767

It's called node port because it creates a virtual port for the node to communicate with. Definition file should have labels bear that's how service knows which pods to communicate with. Use case - you want your particular pod/pods available to access from outside of your network → use nodeport steps → create a definition file - set port no. & values - apply



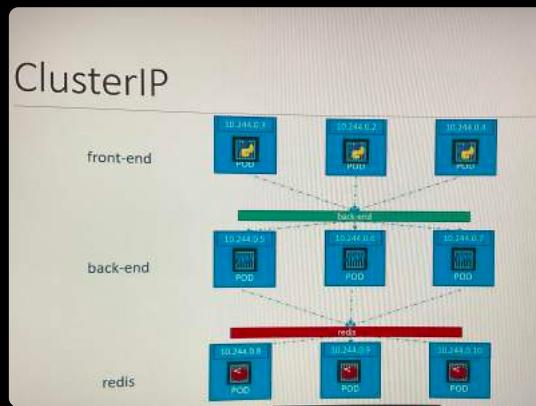
works on
!!!
many
So either that
nodeport is setup on
particular node
we can communicate
with all the pods

Service is highly flexible the same config. steps can be followed for all

② ClusterIP -

Works as a clusterIP for better communication

- central interface to communicate
- Exposes pods within the cluster



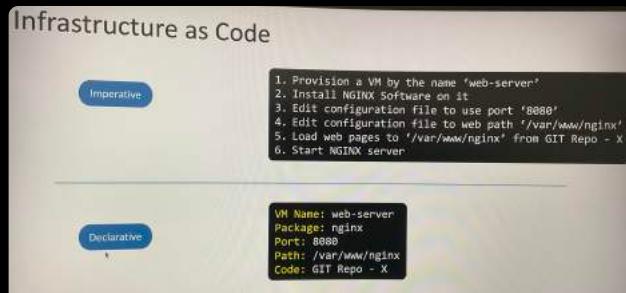
• k get service

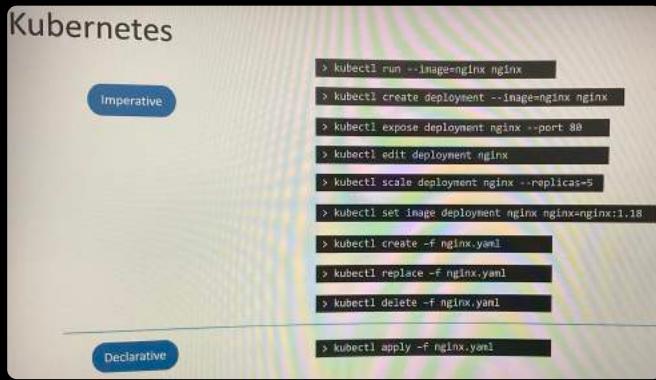
③ Load Balancer -

To make it accessible via single URL.

• kubectl get Service

• ClusterIP is the default service created by kubernetes.



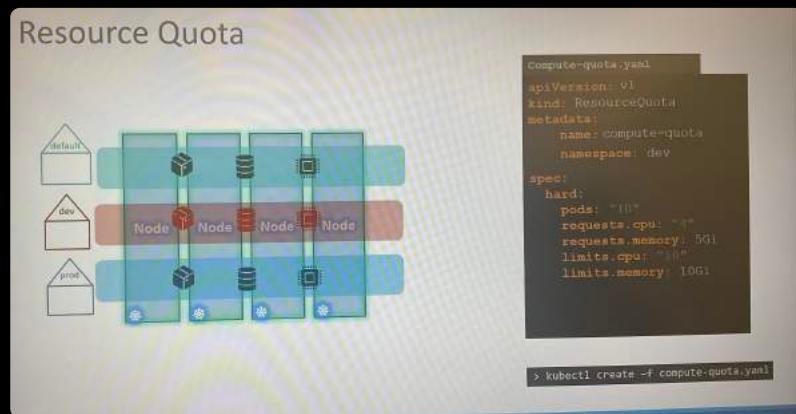


use `kubectl apply`

- Namespaces
- Basically a workspace in terms of terraform
 - we can add namespace value in defn file to make sure pod is operating in right ns.
- To create namespace simply a definition file or
 - `kubectl create namespace nameofns`
- To see resources of dev namespaces
 - `kubectl get pods --namespace=dev`
- To switch namespace
 - `k config set-context $(k config current-context) --namespace=dev`
 - `k get pods --all-namespaces`

• Resource Quota - namespace

To limit resources in namespace we can create resource quota .



- k run pods -----
- k create deployments -----
- k create namespace dev

• Schedulers

- * • IF pod-definition file has nodeName explicitly mentioned that means those nodes are manually scheduled or else if there is not then means these pods are open for auto scheduling
- * If no schedulers then pods are in pending state since the def" file does not have nodeName mentioned.
- * nameNode can only be assign at creation time.
- * To change the nameNode we have to do Binding (creating Binding obj.)
 - To check if scheduler is there
`kubectl get pods --namespace kube-system`
 - If there is no scheduler running the pods will be in the pending state.

To see the available nodes on the cluster

• Labels and selectors

Labels - To filter or group

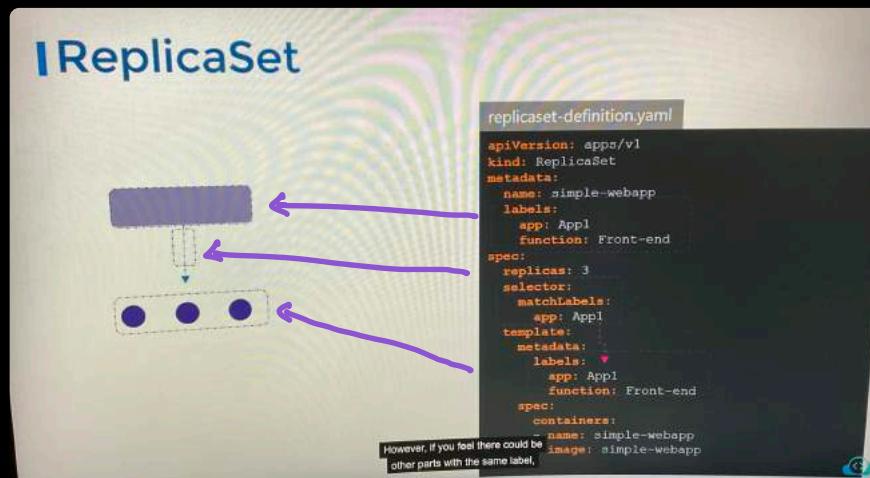
- pod-definition.yaml

Labels:

app: App1

function: front-end

kubectl get pods --selector app=App1



ex.

kubectl get pod --selector env=prod

kubectl get pod --selector bu=finance

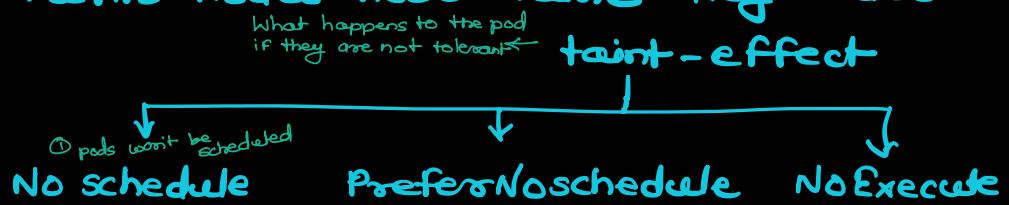
kubectl get all --selector env=prod,tier=frontend

• Taints & Tolerations

Taints and Tolerations are set to create restrictions on the pods to allow or deny pods access to particular nodes.

- Taints are set on nodes
- pods are set on tolerations

- kubectl taint nodes node-name key=value:



ex. kubectl taint nodes node1 app=blue:NoSchedule



- taints and pods does not restrict the pods to go to particular node. Instead it only makes sure that → Particular node only accepts the Tolerated/particular pods.

If goal is to make sure certain pod only go to certain node with filling out limitations of taints and tolerations it's called as node affinity.

- Master node is automatically tainted by default to prevent any pods to run on Master node and that's a best practice.

To see this

```
kubectl describe node kubemaster | grep taint
```

l k get nodes

- kubectl describe node node01 | grep -i taints
- kubectl taint node node1 spray=main:NoSchedule

Reminder -

if yaml file is given
descriptive :-

```
kubectl apply -f pod.yaml
```

Imperative :

```
kubectl run tom --image=nginx
```

To taint node we do this

• `kubectl taint node spray=mortain:NoSchedule`

- for tainting pod we add the tolerations inside the yaml file

ex.

```
GNU nano 4.8
apiVersion: v1
kind: Pod
metadata:
  name: bee
spec:
  containers:
    - image: nginx
      name: bee
  tolerations:
    - key: spray
      value: mortain
      effect: NoSchedule
      operator: Equal
```

- To untaint — to remove taint from node
get the taint from node

1. — `kubectl describe node controlplane`

copy the whole taint from there
and

2. `kubectl taint node controlplane paste -`

Node selectors

- To set particular pod on particular node we can take advantage of node selectors and affinity, but it does not guarantee that this pods won't be placed on other nodes

- To make sure that particular pod go to particular node only and nowhere else we have to use taints and toleration along with node affinity.

Example use-case

- To make sure the high load pods are put inside high size node.
- To basically manage the resources and avoid high load resources to put into smaller node.
- kubectl label nodes node1 size=large.

You cannot say put pod inside a node where node is not small or put pod in Large or med.

To do this we have node affinity

Two ways

- ① Node selectors
- ② Node affinity

Node Affinity ~ To have advance expressions
ex. OR , NOT

- Node affinity :-

Node selector is limited like it does not provide and OR option → to tackle that node affinity which is more Advanced is used.

```
pod-definition.yml
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
spec:
  containers:
    - name: data-processor
      image: data-processor
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
          - matchExpressions:
              - key: size
                operator: In
                values:
                  - Large
```

Node Affinity Types

Available:

`requiredDuringSchedulingIgnoredDuringExecution`
`preferredDuringSchedulingIgnoredDuringExecution`

Planned:

`requiredDuringSchedulingRequiredDuringExecution`
`preferredDuringSchedulingRequiredDuringExecution`

• Resource Request

- We can add "resources" section in pod definition file to set the resources.
- We can also set limits
- If pod constantly try to go above limited resources, Pod will be terminated with out of memory (oom)

```
LimitRange
```

```
limit-range-cpu.yaml
apiVersion: v1
kind: LimitRange
metadata:
  name: cpu-resource-constraint
spec:
  limits:
    - default:
        cpu: 500m
        limit: 500m
      defaultRequest:
        cpu: 500m
        request: 500m
      max:
        cpu: "1"
        limit: 1
      min:
        cpu: 100m
        request: 100m
      type: Container
```

```
limit-range-memory.yaml
apiVersion: v1
kind: LimitRange
metadata:
  name: memory-resource-constraint
spec:
  limits:
    - default:
        memory: 1Gi
      defaultRequest:
        memory: 1Gi
      max:
        memory: 1Gi
      min:
        memory: 500Mi
      type: Container
```

- Resource quota is namespace level object

Editing a Pod

- only certain things can be directly edited in pod
 - spec.tolerations

To edit there are two options

- ① kubectl edit pod podname
then once that changes saved to temp since it gives error when trying to change main file
→ Once done create a new pod by using that file

2. create an yaml file of existing pod

- kubectl get pod webapp -o yaml > pod2.yaml
vi pod2.yaml - To make changes
- then delete old pod and create a new pod

• DaemonSets

- Is one copy always available on each node.
- It is a set which is used to run pod of stack of pods on every single node - ex.
 - for logging nodes or for monitoring we want to make sure every node has this pods with services installed on them running everywhere.

Daemonsets are used for deploying background services on the nodes for ex. system operations service, collecting logs, monitoring frameworks like prometheus and storage volumes

- `kubectl get daemonsets --all-namespaces`
- `kubectl describe daemonset kube-pause --namespace = kube-system`

• Static Pods

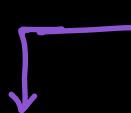
- When we have pods which are created by kublet on its own are known as static pods.

To create such pods copy your yaml files to

`/etc/kubernetes/manifest`

location is saved in `kubelet.service`
2. He can also put config file in `kubelet.service` which consist the config for static pod.

* only pods can be created this way not ds, rs



`k get pods --all-namespaces`

- So with `-control-plane` at the end listed are static pods.

* fun-exercise

— Deleting static node.

1. Look for which node has that pod

`k describe pod -o wide` → node found

2. Get IP of that node since static pods can
only be deleted when we run yaml file

`k describe node node01 -o wide`

3. SSH to that node then find pod location
default → /etc/kubernetes/manifest if not
check `/var/lib/kubelet/config.yaml` get locⁿ
from there

4. Once found run `rm pod.yaml`.

— x —

Custom schedulers

We can have a custom scheduler set.

- We might want to add extra filtering layer to pods before they can be scheduled
- or we want to replace the default scheduler we can do all this

To see currently running scheduler

- `kubectl get events -o wide`.

- `kubectl logs my-custom-scheduler --namespace=kube-system`

- `kube-system` node is where all the system services are located for kubernetes.

- To see active scheduler and other services

- `k get pods --namespace=kube-system`.

- `k get namespaces`

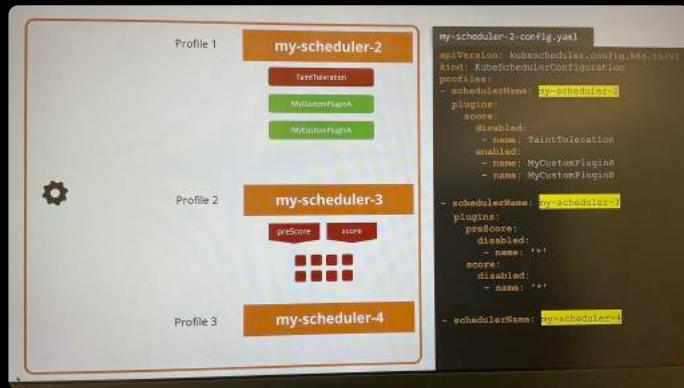
- Scheduling Phase -

1. scheduling Queue
2. filtering
3. Scoring
4. Binding

- We can also set priority to pods to add them in the start of scheduling queue.

- Scheduler profiles

They are used to merge all the custom scheduler at one location to minimize confusion & reduce resource usage and manage all in single binary.



Logging and Monitoring

- In pods services kubelet has a service called **Cloud Advisor** Which sends all the metrics to different metrics services to monitor.

- kubectl top node**
- kubectl top pods**

- Application logs**

- kubectl logs -f podname**

↑

for live logs

- If container contains multiple pods then to get a log for particular pod we have to explicitly mention the name of the pod.
- ex. `k logs -f event-simulator-pod pod-name`
↑ container name

4.0 # • Application Lifecycle management

- Rollouts
 - `kubectl rollout status deployment / myapp-deployment`

- To see history

- `kubectl rollout history deployment / myapp-deployment`

- By default Rolling update

- Where one by one update is done on the system.

Two update strategies

- ① Recreate -

- first it is scaled down to 0 then new version is deployed.

- ② RollingUpdate

- scaled down 1 by 1 & simultaneously scaled up.

- `kubectl rollout undo deployment / myapp-deployment`

- To see the current deployment strategy

- **kubectl describe deployments**

In Kubernetes to run commands on containers which we just created we can use command:
["sleep"] as an entrypoint

4.1 Application Command

- Environmental Variables

* Another way to env is set it directly inside definition file.

1. ConfigMaps -

If we have multiple environment variable in pod-defn file it is better to save them in separate file which is called as configMap

- configMap consist of configuration data

To view configMap.

- **kubectl get configmaps**
- **kubectl describe configmaps**

To view the environmental variables in the pod

- **kubectl describe pod.**

- * As stated before kubectl edit may not allow us to make the required changes

In that case we just save the tmp file and delete old pod and create new or

- kubectl replace --force -f paste location

2. • Secrets

- To store sensitive information

imperative -

```
kubectl create secret generic \  
app-secret --from-literal=DB_Host=mysql  
--from-literal=DB_User=root
```

declarative -

```
kubectl create -f
```

While saving secrets make sure you
Save them in encoded format

↳ To encode `echo -n 'mysql' | base64`

- kubectl describe secrets

- kubectl get secret app-secret -o yaml

↳ To get values.

To decode

```
• echo -n 'encodedtext' | base64 --decode
```

- Secrets are not encrypted they are only encoded.

```
• Secrets are not encrypted in ETCB
```

```
• Anyone able to create pods/ deployments  
in the same namespace can access  
the secrets.
```

- Consider external secrets providers like AWS, GCP.

```
• Anyone with access to etcd can able to see  
encoded information too
```

3. Inside definition file

- You can set env inside definitions file

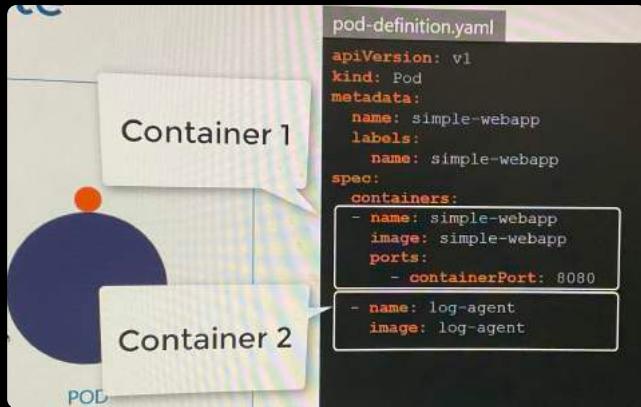
ex. env:

- name: APP_COLOR

valuefrom:

secretkey.

• Multicontainer Pods



- Containers are smallest inside a pod
- To create multi-container setup for the pod change the manifest file → Add two different container names into it

We use multicontainer setups in usually along with 1 pod as main pod and other being log agent



• Cluster Maintenance

If we wanna use the node to upgrade or anything. you can use command

- `kubectl drain node-i`

All pods are gracefully terminated and turned on different node.

and that node-i will be unschedulable until we disable it.

- `kubectl uncordon node-i`

To allow scheduling again

- `kubectl cordon node-i`

Just to make node unschedulable

- no pods termination like drain and

- To empty the node and mark it unsched-

• `k drain node01 --ignore-daemonset`

↑

if we use this command on a pod which is not a part of replicaset it will not work.

- Backup and Restore

two options - querying etcd or backing up file.
If we did some changes imperative way to
document those changes query the kube-apiserver
to save those changes

ex. `kubectl get all --all-namespaces -o yaml > all-deploy-service.yaml`

tools like valero ← for Backup.

- to save a snapshot in ETCD client

`etcdctl snapshot save Snapshot.db`

Different Authentication Methods for kube-apiserver

- ① static file ③ certificates
- ② Tokens ④ Identity services

* Security Basics

- ① Symmetric key - single key
- ② Asymmetric - public & private

5. Security

Three types of certificates

① client

② Root → installed on Root ex. Symantec

③ Server Certificates

Naming convention

Public key

Server.crt

Server.pem

client.pem

Private key

Server.key

Server-key.pem

client.key

*.crt *.*.pem

private key has key mentioned
in their names.

*.key

*-key.pem

- Different Services talks with each other using TLS secure connection
- Every service in k8 has a certificate and private key to communicate among different services.

- To sign this certificate either we can use self sign (not recommended) and signed by CA (certificate Authorities)

To generate certificates

Tools

- ↳ - EsyRSA
- OPENSSL
- CFSSL

View certificates in already existing System.

- `cat /etc/kubernetes/manifests/kube-apiserver.yaml`

it will show all the certificates
once file shows all the certificates
own the

- `openssl x509 -in /etc/kubernetes/pki/ apiserver.crt -text -noout`

- Creating certificates with openssl

① `openssl genrsa -out ca.key 2048`

② `openssl req -new -key ca.key
-subj "/CN=k-CA" -out ca.csr`

③ `openssl x509 -req -in ca.csr -signkey
ca.key -out ca.crt`

④ To view current crt info

`openssl x509 -in file.crt -text -noout`

- certificate API is handled by `Controller Manager`

• `k get csr`

- Adding TLS of new user

- New user joins a team and asking for access to cluster →

① Get CSR and key from user

② encode that csr into base64

`cat newuser.csr | base64 -w 0`

③ Once encoded put the base64 encoded csr to definition file

④ View the status of csr → It will be in the pending state run

`k get csr`

`k certificate approve <name>`

`k describe csr name -o yaml
-o wide`

Inspect Service Log

• `Journalctl -u etcd.service -i`

for kubeadm

`kubectl logs etcd-master`

When kubernetes-api or kubeadm
are down we can still access the
logs using docker commands

`docker ps -a` → To list all pods

To view

• `docker logs 87fc`

To resolve any certificate related issue
which caused any service to stop working
ex. `kubectl is stopped working`

- check docker process logs

• `docker ps -a |grep kube-api`

for CRI compatible resources use

↳ Container runtime interface.

`crictl ps -a`

:2379 → Default port for etcd

`crictl logs containerid from ps`
or

`docker logs containerid`

All certificates related to etcd are saved
at `/etc/kubernetes/pki/etcd`

Certificate API -

To manage certificate signing
process.

• `kubectl get csr`

1. To create a key

• `openssl genrsa -out jane.key 2048`

2. To generate sign in request using that
key

• `openssl req -new -key jane.key
-subj "/CN=jane" -out jane.csr`

To generate Base64

- `Cat filename.csr | base64 -w 0`

To see the status

- `kubectl get csr`

`kubectl certificate approve akshay`
`deny`

kubeconfig

- Provides authentication and authorization information for accessing a k8 cluster.

Assume that I want to access pod info from my cluster (non-admin) in that case I have to pass information like this " k get pods
We cannot put all this info again & again that's why we use kubeconfig
.by default config file is there in Home directory that's why we don't have to pass the info so far.

- `k config view`

kube config default file has three sections
① clusters
② contexts
③ users

```
-- server: my-kube:6443  
-- client-key: admin.key  
-- client-cert: admin.crt  
-- certificate-authority: ca.crt
```

kubeconfig default file path

- `/root/.kube/config`

To know current context —

- `kubectl config --kubeconfig=/root/my-kube-config current-context`

To change current context

```
• kubectl config --kubeconfig=/root/my-kube -  
  config use-context research
```

API

To see API versions -

```
curl https://kube-master:6443/version
```

```
• curl https://kube-master:6443/api/v1/pods
```

Authorization Mechanism

① Node Based

We create a attribute for everything so each policy or authorization will be separate attribute — which will be hard to manage.

② ABAC - Attribute base

— We create roles and that roles can be associated with the particular group.

③ RBAC - Role based

We use webhook and API

④ Webhook -

- Always allow
- always deny

Where do we specify them ?

- they are set on kube-api server

- /etc/kubernetes/manifest/kube-apiserver.yaml

- By default they are set to always allow
can also specify multiple modes

ex - --authorization-mode=Node,RBAC,
webhook

To see the default roles

- k get roles

To see the access of different roles

• k describe ^role kube-proxy -n kube-system

-n = namespace

• k describe rolebinding kube-proxy -n
kube-system

• check access

• k auth can-i delete nodes

• k auth can-i create pods --as dev-user

* How to create user and give permissions

-① create a role → if dev — developer

• kubectl create role developer \

-- verbs=list,create,delete \

-- resource pod

② create rolebindings for that user so
that we can bind that user to
that role

• kubectl create rolebinding dev-user-binding \

-- namespace=default --role=developer \

-- user=Sam

OR

We can just create manifest files

- cluster roles

- Set directly on the cluster so all the nodes inside that cluster follows.

- k get clusterroles

- k get clusterrolebindings

• Service Accounts

Two types of Accounts in kubernetes

- ① User Account
- ② Service Account

• Service Accounts are used by services for monitoring and logging purposes.

ex. prometheus uses service account to get monitoring information

• every namespace has default service account.
• We can use that token to external API services like prometheus.
• In new update you will have to generate token manually

② Service Account

To create

```
kubectl create serviceaccount name
```

To view

```
• kubectl get serviceaccount
```

This creates a token which can be used by services to access the account

To view the token

```
kubectl describe secret token-name
```

This will show the token id

• Image Security

Containers are separated using namespaces
but shares the same kernel

In docker host where all containers are
separated by containers

- If we run `ps aux` command from
docker host. It is also going to show
the processes running in container with
different process id's.

Security context

- Security context provides way to set various aspect of security such as user and group IDs.
- By default Kubernetes is set to all allow
- We can add network policy to the pod which will only allow the traffic which is explicitly allowed.
- To implement the network policy we use labels and selectors.

• Storage

storage in docker - fs → /var/lib/docker

Layered Architecture for docker

- L1 - Base Ubuntu layer

changes in apt packages

changes in pip packages

source code

L5: Update Entrypoint with "flask" cmd.

L6: Container layer - made when
we create container base of this
above layers

volumes -

- docker volume create data-volume.
to mount

docker run -v data-volume:/var/lib/mysql

all volumes are stored in

/var/lib/docker

- Storage drivers are the drivers who takes care of layered Architecture in the docker.

PV - persistent volumes

PVC - claimed by user or pod
way to dynamically request a specific amount and type of storage without needing to know the underlying storage.

Ex. Asking for 50GB storage is called PVC.

- To Bind PV 4 PVC there is Binding criteria
- if no volumes available PVC will be under pending state.

• k get persistentvolumeclaims

• k delete -t

- **pv declaration policy** is set to retain by default
 - means it will stay until it is deleted by the admin.

PV

- A storage in the cluster

PVC

- A request for that storage

• Networking

• Switching :-

Switch is used for local communications.

cmd

• ip link

cmd

• ip addr add 192.168.1.10/24 dev eth0

cmd

• route

Router -

Comm'g bet'w two networks

• route → To display routing table

• To add routing bet'w two networks

• ip route add 192.168.2.0/24 via 192.168.1.1

ip add of diff network where we want to connect

Router's ip address which works as gateway to connect to diff network

• default gateway for all the traffic

• ip route add default via 192.168.1.1

• ip forwarding enable

— cat /proc/sys/net/ipv4/ip-forward

• DNS

• All DNS entry are set inside /etc/hosts file for local but as entry grows it becomes hard to manage such entry that's where DNS server comes into picture.

local host file for entry table

- DNS server keeps track
- Location to set DNS server ip address
 - cat /etc/resolv.conf
- Which file to check when looking for DNS is set in /etc/nsswitch.conf
 - nslookup Sam.com
 - dig Sam.com

- Network Namespaces
- They are used to create a separation betn different network's routing table, ARP tables.
 - so container has no idea about host's Ip address routing table or any other network config.
 - To create network namespaces
 - ip netns add red
 - ip netns
 - To run a ip link command inside red
 - ip netns exec red ip link
 - or
 - ip -n red link

- To see ARP table

- arp

- To add connection bet'n two network namespaces



- Docker Networking

- Below is None network which is limited to that container only — no comⁿ

- docker run --network none nginx

- network option #2

- docker run --network host nginx

- network option #3

• CNI - Container Networking Interface

- set of standards set to develop a networking solution
 - CNI has plugins which use scripts to handle networking problems.
-
- Docker doesn't support CNI directly we can do it with work around.

Ports	
# Name	# Port
• kube-api	6443
• ETCD	2380
• kube-scheduler	10251
• kubelet	10250
• kube-controller-manager	10252

- To get internal IP address of nodes
 - k get nodes -o wide
- To get Mac address of that port
 - ip a | grep -B2 (ip address)
(Make sure you ssh to that node before running)
- To get default gateway
 - ip route show default
 - ifconfig
- How to see ports pod listening on ?
 - netstat -nplt
- CNI in kubernetes
 - CNI location is stored in kubelet.service
 - CNI location - /etc/cni/

* Docker Networking

• `docker info` - To see list of available networks

Types of networks on docker

- Bridge
- Host
- ipvlan
- macvlan
- null
- overlay

By default - Bridge network

① Bridge Network



- With Bridge containers can communicate with each other but inter-bridge communication needs a gateway setup.

② Host Network



- With host Network type Container uses the same Ip address as docker veth
Advantage
 - Due to using same public IP latency reduces and performance increase.

Disadvantage

- Since public IP it's open to internet that means more vulnerable for attacks

We can create multiple containers and all can use same IP address just different ports.

We use service alongwith pods for networking

k get services

k create services

so that service manifest file will have a static ip address assigned to it

Service usage -

- ① Directing traffic to the correct pod
- ② Load Balancing

* Every pod in k8 see's the same IP address

* for labels it is mandatory that both the labels (if you have two labels) They should match exactly.

One matching and other not then that won't be considered as part of that group.

• DNS in Kubernetes

– Kubernetes uses CoreDNS as its DNS provider

– Kubernetes creates a DNS cluster IP for each service

– Kubernetes creates a DNS record for each pod

– Kubernetes creates a DNS record for each node

– Kubernetes creates a DNS record for each namespace

– Kubernetes creates a DNS record for each secret

– Kubernetes creates a DNS record for each config map

– Kubernetes creates a DNS record for each endpoint

– Kubernetes creates a DNS record for each role binding

– Kubernetes creates a DNS record for each role

– Kubernetes creates a DNS record for each service account

– Kubernetes creates a DNS record for each deployment

– Kubernetes creates a DNS record for each stateful set

– Kubernetes creates a DNS record for each daemon set

– Kubernetes creates a DNS record for each config map key

– Kubernetes creates a DNS record for each secret key

– Kubernetes creates a DNS record for each endpoint key

– Kubernetes creates a DNS record for each role binding key

– Kubernetes creates a DNS record for each role key

– Kubernetes creates a DNS record for each service account key

– Kubernetes creates a DNS record for each deployment key

– Kubernetes creates a DNS record for each stateful set key

– Kubernetes creates a DNS record for each daemon set key

– Kubernetes creates a DNS record for each config map key key

– Kubernetes creates a DNS record for each secret key key

– Kubernetes creates a DNS record for each endpoint key key

– Kubernetes creates a DNS record for each role binding key key

– Kubernetes creates a DNS record for each role key key

– Kubernetes creates a DNS record for each service account key key

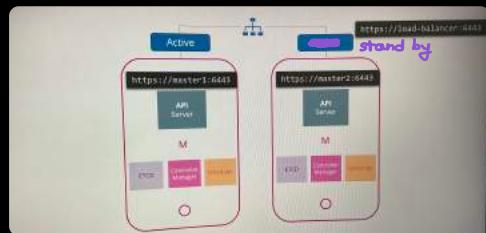
- Design and Install a k8
- Understand the purpose of installing cluster
 - If Education
 - Minikube
 - Single node - AWS
 - Development and Testing
 - Multi-node cluster with a single Master and multiple workers
 - Setup using kubeadm tool or quick provision on GCP or AWS

→ What happens when Master node is not working?

As long as the worker nodes are working functioning of pods is active. However, creation or update cannot be done. We cannot access cluster through API.

• High Availability

- To have high availability its better to use two Master nodes alongwith Load Balancer.



- To prevent any conflicts bet'n two masters and their services it is better to put one Master in Active mode and other onto standby.
- We can decide which Master to put put Active by leader selection process

• kube-controller-manager --leader-elect true

ETCD HA

- When ETCD is installed with same stack as other services its called as stacked topology.

Disadvantage —

When even one master is down etcd becomes unavailable

that's why we can install etcd separately called as **External etcd topology**.

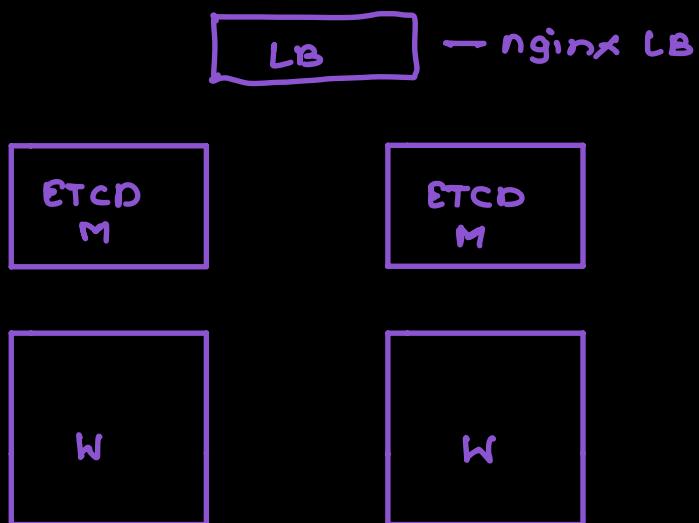
Disadvantage —

- Harder to setup
- requires more numbers of etcd servers.

Where the etcd server is ?

You can check that by

- `cat /etc/systemd/system/kube-apiserver.service`.



Installing kubernetes

- Each node requires container-runtime installed to run containers on the pod for kubernetes
- Docker uses cri-dockerd
- There are different types of containerd available in the market
 - containerd
 - CRI-O

- Troubleshooting in controlplane

if installed by kubeadm

— look for pod in kube-system namespace.

- if installed as service

service kube-apiserver status

kube-controller

kube-scheduler

kubelet

kube-proxy

To check logs here →

• sudo journalctl -u kube-apiserver

check logs of the both

*** kube-system pods are static pods
so their manifest files can be found in
/etc/kubernetes/manifest
Make changes here.

- IF node crashed
 - ① check nodes
 - k get nodes
 - ② k describe node name
 - ③ top
 - ④ df - h
 - ⑤ service kubelet status
 - ⑥ sudo journalctl -u kubelet
- ⑦ Check certificates

• openssl x509 -in /var/lib/kubelet/Worker-1.crt
-text

ex. so node01 was not working properly

To solve this issue

- k get nodes
- describe node
- * — sudo journalctl -u kubelet — to get logs
check reason IF it says stopped
- ssh node01
 - service kubelet status
 - service kubelet start

- Since we are dealing with systemd managed system
we have to use journalctl

- check logs with

Sudo journalctl -u kube-system

ssh to that node all the config related
to kubelet will be at

* * * • /var/lib/kubelet/config.yaml

One more in /etc/kubernetes/kubelet.conf