

## First thing to do is **Authentication**

Once new code added to `terraform (tf)`

First thing we have to do is

1. • `terraform init`

2. • `terraform plan`

To see the tf's upcoming plan

3. `terraform apply`

- The tf file will have the provider section with access key and secret key.

- With `terraform init` it downloads the required dependencies into `.terraform` folder

- `terraform destroy` — To destroy all the resources created in that folder

To destroy specific tf files and resources you can specify the target

• `terraform destroy -target aws_instance.myec2`

- We can also destroy the resource in terraform Just by commenting it out

#### • Terraform State File

- Terraform creates a state file to keep a track of resources

This files are saved with extension `.tfstate`

- The statefile will have more information on the resource other than mentioned info.

#### Desired state & current state

- Desired state - is a state which is desired by the system.
- Current state -

- Terraform's primary function is to create, modify and destroy infrastructure resources to match the desired state described in a TF's configuration.

$v >= 3.0$  means in range of 3.x any.

IF we specify version that info is stored inside `terraform.lock.hcl` file

Once lock file has the HCL we can either delete lock file to change the version or use upgrade to take new constraints.

- `terraform init -upgrade`

Dependency lock file - is used to lock to a specific version of the provider.

- `terraform refresh` is dangerous command try to avoid doing it manually.
- Deprecated in newer versions
  - `-refresh-only` option is always there.

- terraform plan already has terraform refresh in it so no need to run it explicitly

- \* Whenever there is a change which is done manually and which is not mentioned in the terraform explicitly then terraform will have no update over that.

ex. We changed security group from default to custom

But if our terraform file has not mentioned security group then terraform will do nothing about that SG.

**Best Practice** → Try to add as much as information in the terraform file. for granular control.

- Cross-referencing attribute
- Terraform Variables -
- Data Source - Way to fetch already existing resources that were not created by tf itself.

Terraform logging -

We can use attribute **TF\_LOG**

To export it to file

`export TF_LOG_PATH=/tmp/terra.log`

**TRACE** is one of the default logging

Terraform validate -

To validate your code.

- Production tip - try to have separate files for each resource, providers etc.

Dynamic Blocks - To dynamically construct repeatable nested blocks

Terraform taint →

terrafrom apply -replace = "aws\_instance.myec2"

This will destroy and recreate the aws resources again to revert back to the changes where everything was working based on the terraform file.

- In previous versions it was called as terraform taint.

- Splat Expression

It allows us to get list of all attribute

- Terraform Graph

- Terraform graph and we can get that dot file and create a visualization of that

### - Terraform Plan file

`terraform plan -out=path`

↑ filename

- you can save this plan to specific file

and we can use the same "path" file to  
restore the same terraform resources

`terraform apply path`

↑ filename

### • Terraform output

Three ways to get the output

- ① By looking into statefile
- ② By running a output function in tf
- ③ By running `terraform output name` command in CLI

## • Provisioners

Provisioners are used to execute scripts on local or remote machine as part of resource creation or destruction

Ex. Once EC2 is created running application script on it.

### Two types

- Allows us to invoke local executables
  - [ ] ① Local exec
    - ↳ After resource creation
  - [ ] ② Remote exec
    - ↳ Runs on remote resources

Command to ssh to AWS using cmd.

```
ssh -i terraform-key.pem ec2-user@pub-ip.
```

### terraform provisioner

and it should be running inside resource

\* provisioners are not recommended way for remote exec rather than that use EOF

## Terraform Settings

- Dealing with the longer files  
make sure to have separate files.

on Large infrastructure to reduce  
the API calls we can

① We can stop the refresh using

① **-refresh = False Flag**

This will reduce API call

② We can also target

**-target = resource**

• **Zipmap**

- constructs a map from a list of keys  
and corresponding list of values.

Comments in the code - # , // or /\* \*/

# Provisioner types

## ① creation-time provisioner

- only runs during creation, not during updating or any other lifecycle.

If creation-time provisioner fails, the resource is marked as tainted.

## ② Destroy-time provisioner - runs before resource is destroyed

When = destroy

terraform apply -auto-approve

## 6. Terraform Modules & Workspaces

### Acessing child module

- `module.<MODULE NAME>.<OUTPUT NAME>`

Requirements for Publishing Module	
Requirement	Description
GitHub	The module must be on GitHub and must be a public repo. This is only a requirement for the public registry.
Named	Module repositories must use this three-part name format <code>terraform-&lt;PROVIDER&gt;-&lt;NAME&gt;</code>
Repository description	The GitHub repository description is used to populate the short description of the module.
Standard module structure	The module must adhere to the standard module structure.
v.x.y.z tags for releases	The registry uses tags to identify module versions. Release tag names must be a semantic version, which can optionally be prefixed with a v. For example, v1.0.4 and 0.9.2



### Terraform Workspace -

- just like windows workspace for separation of work.

## 7. Remote state management

.gitignore  
files to ignore  
.terraform  
terraform.tfvars → may have sensitive info  
terraform.tfstate →  
crash.log → log file

Terraform and .gitignore	
Depending on the environments, it is recommended to avoid committing certain files to GIT.	
Files to Ignore	Description
.terraform	This file will be recreated when terraform init is run.
terraform.tfvars	Likely to contain sensitive data like usernames/passwords and secrets.
terraform.tfstate	Should be stored in the remote side.
crash.log	If terraform crashes, the logs are stored to a file named crash.log

Terraform state file is being used by a team so storing it locally is not a good option instead of that we can store them on the Backend server so everyone can access them.

But that file should have locking mechanism to protect the file being changed by multiple users at a time which may create inconsistency.

# Example of Central Backend

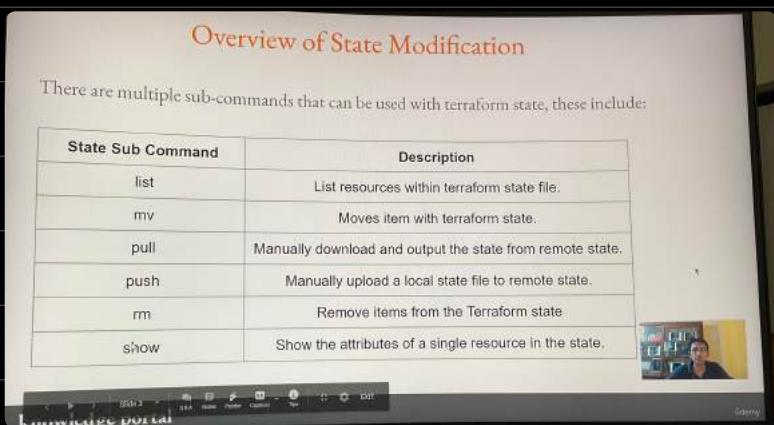
## - AWS S3 Bucket

State locking - Whenever we are performing terraform operation then state file will be locked to protect from corruption of file.

\* Not All Backends supports state locking.

- force-unlock → to manually unlock
- S3 does not support the state locking
- With the help of dynamodb we can state lock

## • Terraform state Management



## • Terraform Import

- With this we can manually import Architecture into terraform.

To do that we create the same tf file as the config we have running and map that to the tf file.

## 8. Security Primer

- To have multiple providers for different resources we create alias in terraform and we can put the new provider block in the resource with the alias to run that resource in that alias.

## • Terraform Cloud

for running TF on cloud

- can add variables
- secure TF state file
- Auto plans on changes
- & many more

## • Sentinels

- Policy as code.

ex. Block any EC2 with no tags  
Block All traffic sg's

\* When doing Remote execution in the Terraform plan the statefile and other files are stored inside a remote cloud while logs are stored locally.

**provisioners** - They allows you to execute scripts or commands on resource after it has been created or updated.

primarily used for task such as bootstrapping, config management.

- ① **remote-exec**

ex. Accessing EC2 instance using TF and running nginx command on remote m/c

② **local-exec** - To run scripts locally.

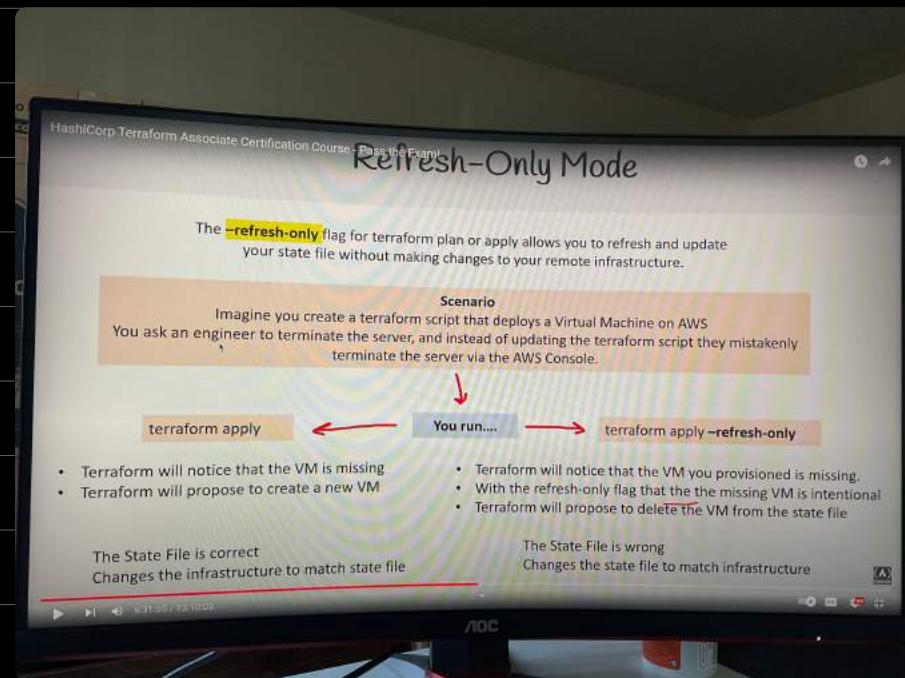
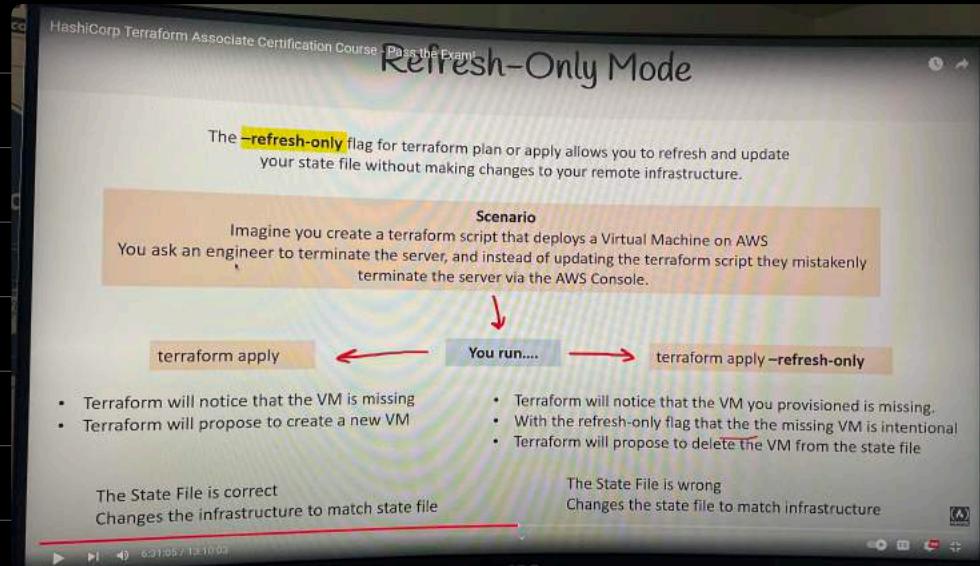
? can modules be stored on github

modules versioning providers -

- not on public registry.

## Notes

- state are saved in `terraform.tfstate`
  - All terraform state subcommands that modify state will write a backup file.
    - Terraform will take the current state and store it in a file called `terraform.tfstate.backup`.
  - Backups cannot be disabled.  
To get rid of the backup file you need to manually delete the file.
- `terraform init`
- Downloading plugin dependencies
    - create a `.terraform` directory
    - create a dependency lock file to enforce expected version to plugin and tf.
  - first cmd to run
  - modify or changed dependencies? → run `tf init`



**Terraform init will create following**

- **.terraform directory**
- **.terraform.lock.hcl file**

- Does terraform plan creates an state file?
- - locks an state file?

TF validate works when tf plan & tf apply

- TF init is required to download tf module
- TF validate doesn't catch any errors related to value ex. ami = "Sam"
- tf providers to get the list of current providers
- Modules can be stored into your private registry.
- By default you are using the backend state
- tf cloud requires you to have workspace.
- When running tf cloud credentials needs to be configured on the cloud environmental variable block.
- We can set partial settings for the Backends using the -backend-config flag for terraform init

ex. `terraform init -backend-config = backend.hcl`

- Terraform remote state allows that state to be stored on the cloud.
- You can disable lock with `-lock flag`
- `tf` When using `save planned` it will not prompt you to confirm and will act like `auto-approve`
- `tf apply` is also running `tf validate`
- Migrating `tf Backend`
- difference between public modules and private modules
- public modules format

`<Namespace>/<NAME>/<PROVIDER>`

ex. `hashicorp/consul/aws`

private modules

- `<HOSTNAME>/<NAMESPACE>/<NAME>/<PROVIDER>`

ex. [app.terraform.io/example/vpc/aws](https://app.terraform.io/example/vpc/aws).

publishing Modules follows -

- versioning
- Readme's
- Well Documentation
- Show examples
- automatically generated documentation

• public modules are managed via public git repo on Github.

It should follow all the tf rules.

• verified modules are reviewed by Hashicorp but managed by users.

Modules directory consist of Nested modules

• private module registry supports

- module versioning
- searchable and filterable list
- a configuration designer

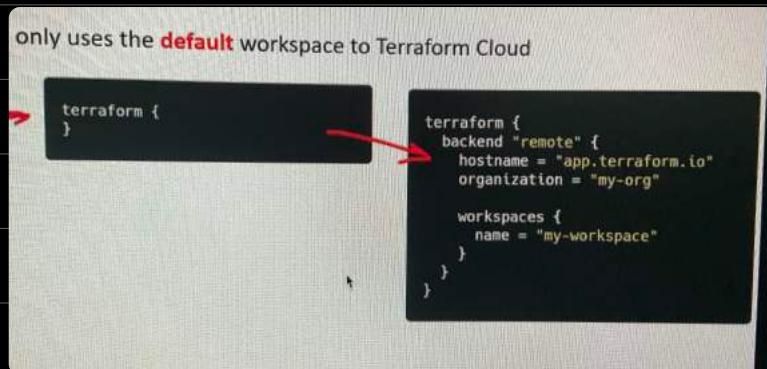
- "cost estimation" feature is only available for teams and Governance plan and above.

- **Workspaces**

- You can choose any tf versions for workspaces
- You can share state globally.
- You can choose to auto approve runs.

## Migrating from default workspace to tf cloud

- create a workspace in tf cloud
- Replace your terraform configuration with a remote backend



- Run `terraform init`

## Workspaces

- CLI Workspaces
- Terraform cloud workspaces

`tf show -json` → to get output in json

- credentials should be stored in
  - ① Environmental variables
  - ② credentials file

- Modules are stored in .terraform/ modules

- Source argument is necessary to call child modules

Modules can use following meta-arguments

- for-each
- depends-on
- count
- provider

- Two complex types in terraform

- ① A collection type
- ② A structural type.

## ① Collection types

- ① List -
- ② set -
- ③ Map -

## ② structural

- Tuples
- objects
- resources
- modules

- Whenever configuration Backend changes you must run `terraform init`
  - default backend - local
  - You cannot load additional backends as plugins.
  - You don't need to configure backend for tf cloud.
- \* • `terraform` / `terraform.tfstate` file contains the backend configurations
- \* partial configuration is allowed
- \* When new backend is detected by `terraform` at the time of reinitialization
- \* To ensure correct operation & destroy `terraform` retains copy of the most recent set of dependencies within the state.
- \* Tf stores metadata.

- \* Tf state increases performance
- State file's are cloud agnostic - means irrelevant to particular cloud provider

- Refreshing state can be made optional for `state` file or can use `-target`

\*

## Terraform cloud.

- Workspaces in terraform cloud and terraform CLI works differently
  - Workspaces are required in tf cloud.
  - Terraform plan will write data to the `terraform.tfstate` file.
  - Terraform refresh — to refresh state file to actually align with real-world
- ① single Workspace can have multiple state files.

private registry modules has Hostname mentioned

- Syntax — `<Hostname>/<NAMESPACE>/<NAME>/<PROVIDER>`

Terraform is immutable, declarative IAC

providers plugins downloaded are stored inside  
• .terraform/providers

• Terraform workspace select dev ← To switch up

• Github is not a supported backend type.

• invalid is a valid variable name in terraform

• Terraform open source stores local state at

• .terraform.tfstate.d/<workspace name>

• cloud agnostic — can run on any cloud platform

• Terraform uses a configuration file to understand resource order and if it is missing then order cannot be determined with config. alone.  
so it stores that into State file.

• You can tf code without providers block. It is not mandatory to have terraform block.

• Variables declared inside modules cannot be called directly

→ you can call variables outside of module directly with the help of output block type

- OSS stores its state file in local backend
- Better way to secure the configuration
  - credential file
  - environmental variable
- To set backend in terraform we can use`terraform {  
 backend "name" {}  
}`
- We can only configure one backend block.
- You cannot install terraform providers from source code.

